

Phase 2 - MLflow Experiment Tracking and Model Registry

Project context

This phase extends the Phase 1 fuel-price regression workflow with local MLflow tracking. The dataset used in all experiments is ``global_fuel_prices_2020_2026.csv``, and the prediction target is ``petrol_usd_liter``.

Objective

The goal of Phase 2 is to:

- ? track multiple regression experiments in MLflow
- ? log parameters, metrics, tags, and artifacts
- ? compare model performance
- ? register the best-performing model in the MLflow Model Registry
- ? compute confidence intervals as a bonus step with ``statsmodels``

Dataset and preprocessing

The preprocessing logic is reused from Phase 1 through ``pipeline.py``. The notebook applies the same core steps:

- ? automatic date feature extraction
- ? dropping leakage-related target alternatives
- ? numeric imputation and scaling
- ? one-hot encoding for categorical variables
- ? identical train/test split strategy

This guarantees consistency between Phase 1, Phase 2, and the later Phase 3 dashboard.

Tracked models

The Phase 2 notebook compares five configurations:

1. LinearRegression
2. Ridge with $\alpha = 0.5$
3. Ridge with $\alpha = 5.0$
4. Lasso with $\alpha = 0.001$
5. RandomForestRegressor with 300 trees and max depth 12

The notebook evaluates all runs under the same preprocessing pipeline.

MLflow logging

Logged parameters

Examples of logged parameters:

- ? `test\_size`
- ? `random\_state`
- ? number of numeric features
- ? number of categorical features
- ? model-specific hyperparameters such as `alpha`, `n\_estimators`, or `max\_depth`

Logged metrics

The notebook logs at least the following evaluation metrics for every run:

- ? `MAE`
- ? `RMSE`
- ? `R2`

Logged tags

The following tags are explicitly written:

- ? `algorithm`
- ? `dataset\_version`
- ? `target`

Logged artifacts

The notebook logs multiple artifacts, including:

- ? residual distribution plots
- ? actual-vs-predicted plots
- ? `conf\_intervals.csv`
- ? `prediction\_ci.csv`
- ? confidence interval visualizations

Bonus: confidence intervals

As requested in the assignment, the notebook fits a `statsmodels` OLS model on the scaled numeric features and computes 95 percent confidence intervals.

Generated files include:

- ? `conf\_intervals.csv`
- ? `prediction\_ci.csv`

These artifacts are intended for later analysis and for reuse in Phase 3.

Best model and registry

According to `model\_meta.json`, the best model selected in Phase 2 is:

? `RandomForest\_300trees`

The notebook then registers the best model in the MLflow Model Registry using:

? `BestFuelRegressor\_v1`

and transitions the newest version to:

? `Production`

Result summary

Phase 2 is largely implemented and fulfills the coding requirements of the assignment:

- ? local MLflow experiment setup
- ? multiple tracked runs
- ? parameter logging
- ? metric logging
- ? artifact logging
- ? tag logging
- ? model registration
- ? production-stage transition
- ? bonus confidence interval logic

Submission status

The following deliverables are present in code or generated-output form:

- ? `phase2.ipynb`
- ? `pipeline.py`
- ? `mlflow.db`
- ? `conf\_intervals.csv`
- ? `prediction\_ci.csv`
- ? `model\_meta.json`

The following parts are still manual submission tasks:

- ? `phase2.pdf` screenshots of the MLflow UI
- ? insertion of explicit screenshots if required by the teacher

Screenshot placeholders

Insert the following screenshots into the final submission if needed:

1. MLflow experiment overview page
2. run comparison table
3. registered model page
4. production model version page

Conclusion

Phase 2 successfully extends the Phase 1 solution into an MLOps workflow with experiment tracking, model comparison, model registration, and uncertainty estimation. The notebook and generated artifacts provide a solid basis for the final Streamlit dashboard in Phase 3.